

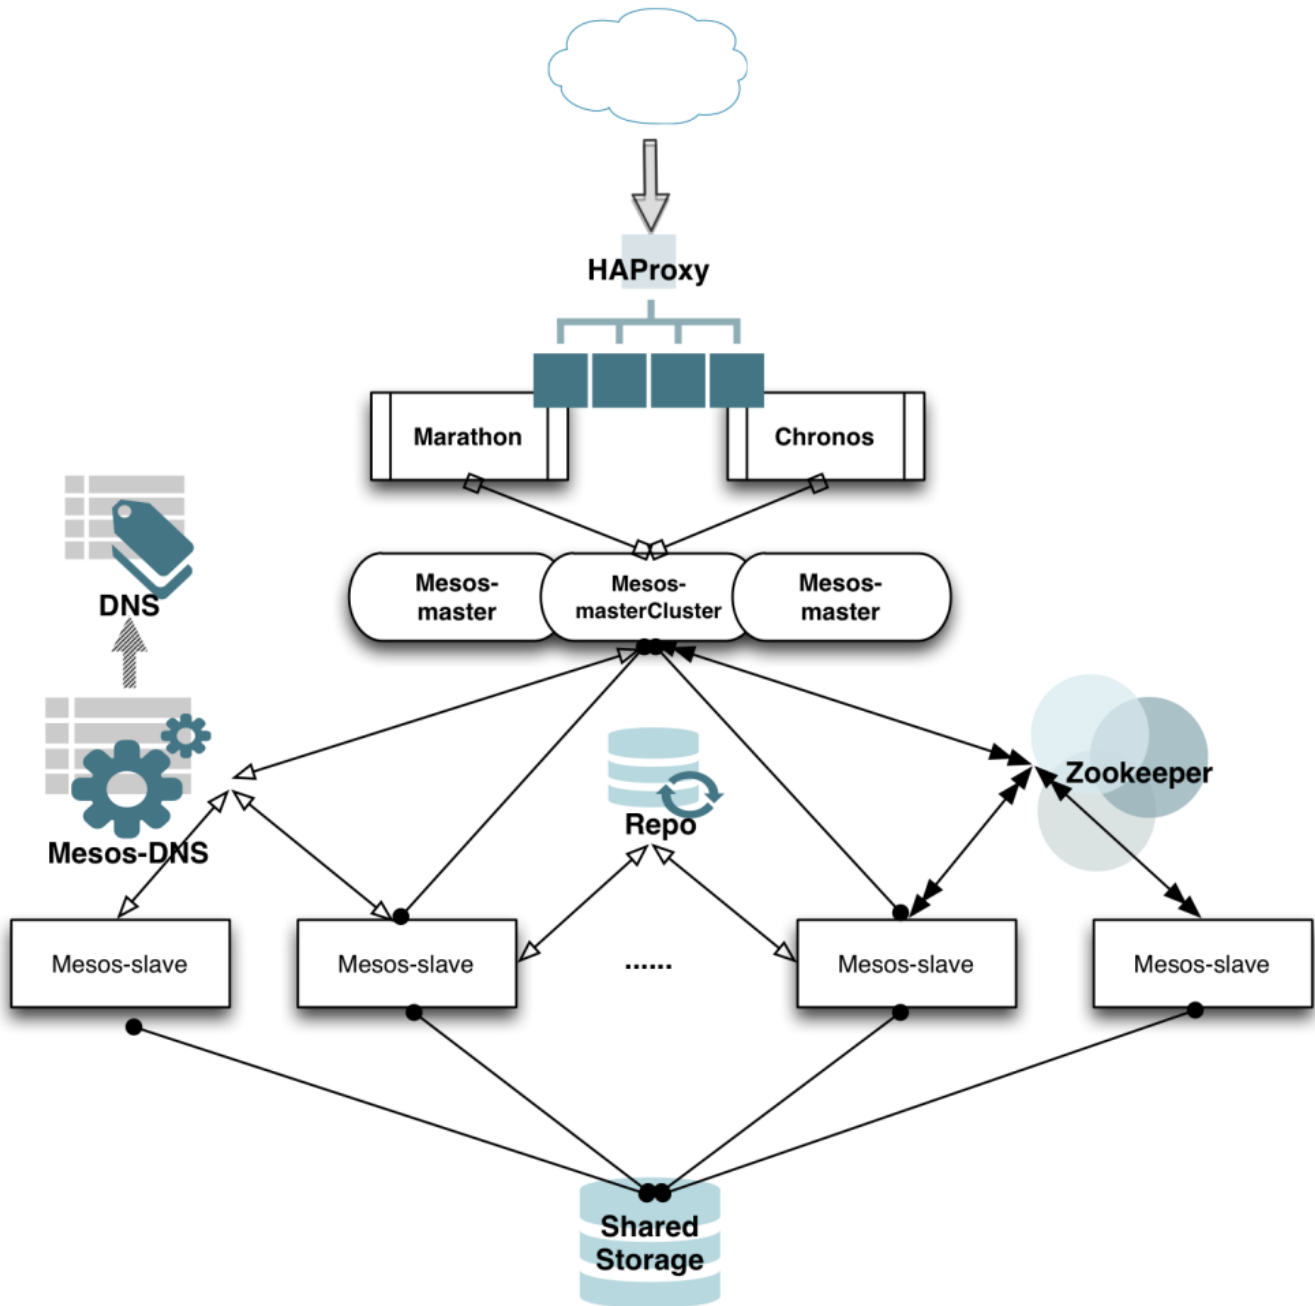
Container Platform With Mesos

Summary:

While PoC is still underway, we would like to present one of the proposed platform architects for deploying Dockerized services on Mesos frameworks. A few services including Kafka, Zookeeper and RabbitMQ have been successfully deployed on such platform.

Reference Architect:

Using Marathon and Chronos frameworks, application/service deployment to various environments become much simpler comparing with CF based platform. Both containerized and regular applications can be seamlessly deployed on Mesos abstraction platforms without complex configuration or run-time environment modification. Leveraging service discovery automation and load balancing functions, multiple service instances can be easily deployed to the platform with full redundancy and without manual intervention.



Considerations and Benefits:

Stability: The approach should be simple and repeatable. The platform should be stable and supportable by engineers from GE.

Fault Tolerance: The platform should have no single point of failure on service. All controller services are clustered without single point of failure. Failed services will be restarted automatically. Services and their dependencies are isolated by containers.

Security: All customer facing API are secured with SSL and password protection.

Agility: Service can be dynamically deployed and migrated to another host. Seamless upgrades are norm.

Capacity Management: Addition nodes can be dynamically added or removed from the mesos-slave farm.

Infrastructure Agnostic: The reference architect can be deployed on AWS, Bare Metals and Virtual servers.

Recommended Procedures and Best Practice:

1. Build a Docker container image with intended application. Carefully craft the application configuration and entry point startup script to handle multiple instances and clustering deployments.

Example: A Dockerfile for a 3-node RabbitMQ cluster

```
alfreds@sjc4dockerdev01:~/rabbitmq$ cat Dockerfile
FROM ubuntu:trusty

# add our user and group first to make sure their IDs get assigned consistently, regardless of whatever dependencies get added
RUN groupadd -r rabbitmq && useradd -r -d /var/lib/rabbitmq -m -g rabbitmq rabbitmq
ENV RABBITMQ_VERSION 3.5.6-1
ENV RABBITMQ_LOGS=- RABBITMQ_SASL_LOGS=-
ENV PATH /usr/lib/rabbitmq/bin:$PATH

# grab gosu for easy step-down from root
RUN apt-key adv --keyserver-options http-proxy=http://3.39.89.55:8080 --keyserver-options https-proxy=http://3.39.89.55:8080 --keyserver ha.pool.sks-keyservers.net --recv-keys B42F6819007F00F88E364FD4036A9C25BF357DD4 \
    && apt-key adv --keyserver-options http-proxy=http://3.39.89.55:8080 --keyserver-options https-proxy=http://3.39.89.55:8080 --keyserver ha.pool.sks-keyservers.net --recv-keys 434975BD900CCBE4F7EE1B1ED208507CA14F4FCA \
    && apt-key adv --keyserver-options http-proxy=http://3.39.89.55:8080 --keyserver-options https-proxy=http://3.39.89.55:8080 --keyserver ha.pool.sks-keyservers.net --recv-keys F78372A06FF50C80464FC1B4F7B8CEA6056E8E56

RUN export http_proxy=http://3.39.89.55:8080 \
    && export https_proxy=http://3.39.89.55:8080 \
    && apt-get update && apt-get install -y curl ca-certificates --no-install-recommends && rm -rf /var/lib/apt/lists/* \
    && curl -o /usr/local/bin/gosu -SL "https://github.com/tianon/gosu/releases/download/1.3/gosu-${dpkg --print-architecture}" \
    && curl -o /usr/local/bin/gosu.asc -SL "https://github.com/tianon/gosu/releases/download/1.3/gosu-${dpkg --print-architecture}.asc" \
    && gpg --verify /usr/local/bin/gosu.asc \
    && rm /usr/local/bin/gosu.asc \
    && chmod +x /usr/local/bin/gosu \
    && echo 'deb http://www.rabbitmq.com/debian testing main' > /etc/apt/sources.list.d/rabbitmq.list \
    && echo 'deb http://packages.erlang-solutions.com/debian jessie contrib' > /etc/apt/sources.list.d/erlang.list \
    && apt-get update && apt-get install -y rabbitmq-server=${RABBITMQ_VERSION} --no-install-recommends && rm -rf /var/lib/apt/lists/*

RUN echo "[{rabbit, [{tcp_listeners, [RABBITMQ_PORT]}, {loopback_users, []}, {collect_statistics_interval, 10000}, {heartbeat, 10}, {channel_max, 5000}, {cluster_partition_handling, autoheal}, {cluster_nodes, [{RABBITMQ_CLUSTER_NODES}], {disc}}, {rabbitmq_management, [{http_log_dir, '/tmp/rabbit-mgmt'}, {listener, [{port, RABBITMQ_MANAGEMENT_PORT}]}]}, {vm_memory_high_watermark, 0.4}, {disk_free_limit, 100000000}, {log_levels, [{connection, info}, {mirroring, info}]}, {delegate_count, 32}, {tcp_listen_options, [binary, {packet, raw}, {reuseaddr, true}, {backlog, 128}, {nodelay, true}, {exit_on_close, false}, {keepalive, true}]}, {collect_statistics_interval, 60000}, {rabbitmq_management_agent, [{force_fine_statistics, true}]}], {kernel, [{net_ticktime, 30}]}]." > /etc/rabbitmq/rabbitmq.config \
    && echo "rabbitmq-server -detached; rabbitmqctl stop_app; rabbitmqctl join_cluster rabbit@RABBITMQ_SEED_NODE || true; rabbitmqctl start_app; rabbitmqctl set_policy cluster-all-queues '^(?!amq\\.)*' '{\"ha-mode\":\"all\", \"ha-sync-mode\":\"automatic\"}'" > /usr/lib/rabbitmq/bin/rabbitmq-member \
    && chmod +x /usr/lib/rabbitmq/bin/rabbitmq-member

VOLUME /var/lib/rabbitmq

# add a symlink to the .erlang.cookie in /root so we can "docker exec rabbitmqctl ..."
RUN ln -sf /var/lib/rabbitmq/.erlang.cookie /root/

COPY docker-entrypoint.sh /
RUN chmod +x /docker-entrypoint.sh
ENTRYPOINT ["/docker-entrypoint.sh"]

EXPOSE 5672 4369 15672
CMD ["rabbitmq-server"]
alfreds@sjc4dockerdev01:~/rabbitmq$
```

Example: A docker entry startup file.

```
alfreds@sjc4dockerdev01:~/rabbitmq$ cat docker-entrypoint.sh
#!/bin/bash
set -e
[[ -z "${RABBITMQ_PORT}" ]] && RABBITMQ_PORT=5672
[[ -z "${RABBITMQ_SEED_NODE}" ]] && RABBITMQ_SEED_NODE='localhost'
[[ -z "${RABBITMQ_MANAGEMENT_PORT}" ]] && RABBITMQ_MANAGEMENT_PORT=15672
[[ -z "${RABBITMQ_CLUSTER_NODES}" ]] && RABBITMQ_CLUSTER_NODES='rabbit@localhost'
[[ -z "${RABBITMQ_ERLANG_COOKIE}" ]] && RABBITMQ_ERLANG_COOKIE="MpRI2iDWBGJ6Y0Z23q9VNHeK"

if [ "${RABBITMQ_ERLANG_COOKIE}" ]; then
    cookieFile='/var/lib/rabbitmq/.erlang.cookie'
    echo "${RABBITMQ_ERLANG_COOKIE}" > "$cookieFile"
    chmod 600 "$cookieFile"
    chown rabbitmq "$cookieFile"
fi

sed -i "s/RABBITMQ_PORT/${RABBITMQ_PORT}/" /etc/rabbitmq/rabbitmq.config
sed -i "s/RABBITMQ_MANAGEMENT_PORT/${RABBITMQ_MANAGEMENT_PORT}/" /etc/rabbitmq/rabbitmq.config
sed -i "s/RABBITMQ_CLUSTER_NODES/${RABBITMQ_CLUSTER_NODES}/" /etc/rabbitmq/rabbitmq.config
sed -i "s/RABBITMQ_SEED_NODE/${RABBITMQ_SEED_NODE}/" /usr/lib/rabbitmq/bin/rabbitmq-member
if [ "${RABBITMQ_SEED}" ]; then
    exec "$@"
else
    myprog="/usr/lib/rabbitmq/bin/rabbitmq-member"
    exec "$myprog" "$@"
fi
alfreds@sjc4dockerdev01:~/rabbitmq$
```

2. Test the container image and push it to the local repository after done.

Example: Login to the local repo and push the container image to repo

```
# Display local Docker images
root@sjc4dockerdev01:~# docker images
REPOSITORY          TAG          IMAGE ID          CREATED          VIRTUAL SIZE
rabbitmq-c31        latest      3f0647ea738b     23 hours ago    374.3 MB

## Log in to the local Docker repository
docker login -u lab1 -p lab1 -e lab1@ge.com https://repo.vms.crd.ge.com:8080

# Tag the target image
#docker tag 3f0647ea738b repo.vms.crd.ge.com:8080/rabbitmq-c3:latest

# Push the image to the local repository
#docker push repo.vms.crd.ge.com:8080/rabbitmq-c3:latest

#Verify the image is on the local repository
#Docker search repo.vms.crd.ge.com:8080/rabbitmq
NAME                DESCRIPTION    STARS     OFFICIAL    AUTOMATED
library/rabbitmq-c3    0
```

3. Generate a json file based Marathon API v2. Provide environment parameters, port mapping details and other application specific configuration values.

Example: A json file to deploy rabbitmq via Marathon API

```
SF01212464633A:marathon 212464633$ cat rabbitmq-c3-0.json
{
  "container": {
    "type": "DOCKER",
    "docker": {
      "image": "repo.vms.crd.ge.com:8080/rabbitmq-c3:latest",
      "network": "BRIDGE",
      "hostname": "rcnode1",
      "name": "rc1",
      "portMappings": [
        { "containerPort": 14369, "hostPort": 0, "servicePort": 14369, "protocol": "tcp" },
        { "containerPort": 5672, "hostPort": 0, "servicePort": 5772, "protocol": "tcp" },
        { "containerPort": 15672, "hostPort": 0, "servicePort": 15772, "protocol": "tcp" }
      ]
    }
  },
  "id": "rabbitmq-c3-0",
  "instances": 1,
  "constraints": [["hostname", "UNIQUE"]],
  "env": { "RABBITMQ_SEED": "1", "RABBITMQ_CLUSTER_NODES": "'rabbit@rcnode1','rabbit@rcnode2','rabbit@rcnode3'", "ERL_EPMD_PORT": "14369" },
  "cpus": 0.5,
  "mem": 512,
  "uris": [
    "file:///etc/docker.tar.gz"
  ]
}
```

In this example, the rabbitmq service port is listening on <vip>:5772 and admin port is on <vip>:15572. The VIP is the virtual IP address of the load balancer. In case of clustering, the first node should use 0.0.0.0:14369 for EPMD service and sub-sequential nodes should connect to port <vip>:14369.

It's highly recommended to group multiple service into one json for a clustering service, particularly with dependency. Here is a generic example for illustrate how it works.

```
SF01212464633A:marathon 212464633$ cat prods2.json
{
  "id": "/prods2",
  "apps": [
    {
      "id": "service-common",
      "cpus": 0.1,
      "mem": 1,
      "ports": [0],
      "cmd": "python -m SimpleHTTPServer $PORT0",
      "instances": 6
    }
  ],
  "groups": [
    {
      "id": "/prods2/task1",
      "apps": [
        { "id": "/prods2/task1/job1", "constraints": [["hostname", "UNIQUE"]], "instances": 6, "cpus": 0.1, "mem": 1, "cmd": "for i in {1..3} ; do echo `date`: 'Hello Marathon/DCOS from job1' >> /tmp/job1.out; sleep 30 ; done"},
        { "id": "/prods2/task1/job2", "instances": 6, "cpus": 0.1, "mem": 1, "cmd": "for i in {1..3} ; do echo `date`: 'Hello Marathon/DCOS from job2' >> /tmp/job2.out; sleep 30 ; done"}
      ]
    }, {
      "id": "/prods2/task2",
      "dependencies": ["/prods2/task1"],
      "apps": [
        { "id": "/prods2/task2/job11", "constraints": [["hostname", "UNIQUE"]], "instances": 6, "cpus": 0.1, "mem": 1, "cmd": "[[ -f /tmp/job1.out ]] && echo `date`: `tail -1 /tmp/job1.out` >> /tmp/job11.out`"},
        { "id": "/prods2/task2/job22", "instances": 6, "cpus": 0.1, "mem": 1, "cmd": "[[ -f /tmp/job2.out ]] && echo `date`: `tail -1 /tmp/job2.out` >> /tmp/job22.out`"}
      ]
    }
  ]
}
```

4. Call appropriate Marathon API to deploy the service. Obtain service entry URL and credentials if needed.

Example: To deploy a json file and discover the deployed service.

```
#!/bin/bash
function usage() {
cat <<EOF
usage: $0 options

Deploy applications on LAB using Marathon

Example:
    $0 -[l|a|d|h|g|G|D] [<app_name>]

OPTIONS:
    -h -- Display help
    -l -- List deployed applications
    -a -- Add a new application
    -d -- Delete an existing application
    -g -- List a group
    -G -- Add a new group
    -D -- Delete a group

EOF
}
while getopts "ha:lgd:G:D:" OPTION; do
case "$OPTION" in
l)
    curl http://adminUser:topP1ssw0rd@3.39.90.170:8080/v2/apps | python -m json.tool
    ;;
h)
    usage
    exit 0
    ;;
a)
    curl -X POST http://adminUser:topP1ssw0rd@3.39.90.170:8080/v2/apps -d @$2.json -H "Content-type:application/json"
    ;;
G)
    curl -X POST http://adminUser:topP1ssw0rd@3.39.90.170:8080/v2/groups -d @$2.json -H "Content-type:application/json"
    ;;
g)
    [[ $# == 1 ]] && curl -X GET http://adminUser:topP1ssw0rd@3.39.90.170:8080/v2/groups | python -m json.tool
    [[ $# == 2 ]] && curl -X GET http://adminUser:topP1ssw0rd@3.39.90.170:8080/v2/groups/$2 | python -m json.tool
    ;;
d)
    curl -X DELETE http://adminUser:topP1ssw0rd@3.39.90.170:8080/v2/apps/$2
    ;;
D)
    curl -X DELETE http://adminUser:topP1ssw0rd@3.39.90.170:8080/v2/groups/$2
    ;;
\?)
    echo "Invalid option: -$OPTARG" >&2
    usage
    exit 1
    ;;
:)
    usage
    exit 1
    ;;
esac
done
```

...and to publish and subscribe from the deployed rabbitmq. All REST APIs will be SSL enabled and protected by password.

```
#./lab.sh -a rabbitmq-c3-0
```

Here are a few running services in the Lab ecosystem.

ID	Memory (MB)	CPUs	Tasks / Instances	Health	Status
/basic-0	30	0.6	5 / 6	<div></div>	Waiting
/basic-1	32	0.4	1 / 1	<div></div>	Running
/basic-3	32	0.5	1 / 1	<div></div>	Running
/basic-busybox	20	0.5	5 / 5	<div></div>	Running
/department-a/product-a/service-a	64	0.2	2 / 2	<div></div>	Running
/department-a/product-a/service-b	64	0.2	2 / 2	<div></div>	Running
/department-a/service-common	64	0.2	2 / 2	<div></div>	Running
/kafka-1	512	0.5	1 / 1	<div></div>	Running
/kafka-2	512	0.5	1 / 1	<div></div>	Running
/kafka-3	512	0.5	1 / 1	<div></div>	Running
/memcached	192	1.5	3 / 3	<div></div>	Running
/rabbitmq-03	1536	1.5	3 / 3	<div></div>	Running
/rabbitmq-c3-0	512	0.5	1 / 1	<div></div>	Running
/rb1	1536	1.5	3 / 3	<div></div>	Running
/redis-master	512	0.5	1 / 1	<div></div>	Running
/redis-slave	1536	1.5	3 / 3	<div></div>	Running
/zk-1	512	0.25	1 / 1	<div></div>	Running
/zk-2	512	0.25	1 / 1	<div></div>	Running
/zk-3	512	0.25	1 / 1	<div></div>	Running

5. Verify and test accessibility and performance of the deployed service.

Example: To verify the deployed rabbitmq service. Test scripts are application specific and need to be written accordingly.

```

root@ctl3:~/bin# cat hello_pub.py
#!/usr/bin/python
from kombu import Connection
import datetime, getopt, sys

hostname='127.0.0.1'
port=5672
username='guest'
password='guest'
queue='test_queue'
message='hello world'

def usage(message=None):
    print "Usage: %s [-h] [-u|--username <user_name>] [-w|--password <password>] [-h|--host <rabbitmq_hostname>] [-p|--port <service_port>] [-q|--queue <queue_name>] [-m|--message <message>]" % (sys.argv[0])
    print "\t-h|--help: show this message"
    print "\t-v|--verbose: include details in output"
    print "\t-u|--username: Username"
    print "\t-w|--password: User password"
    print "\t-h|--hostname: Rabbitmq hostname or IP"
    print "\t-p|--port: Rabbitmq service port"
    print "\t-q|--queue: Rabbitmq queue name"
    print "\t-m|--message: Message to be published"
    sys.exit(-1)

(opts, args) = getopt.getopt(sys.argv[1:], "?vu:wh:p:q:m:", ["Rabbitmq test", "help", "verbose", "username", "password", "host", "port", "queue", "message"])

for o, a in opts:
    if o in ["-?", "--help"]:
        usage()
    elif o in ["-u", "--username"]:
        username=a
    elif o in ["-w", "--password"]:
        password=a
    elif o in ["-h", "--hostname"]:
        hostname=a
    elif o in ["-p", "--port"]:
        port=a
    elif o in ["-q", "--queue"]:
        queue=a
    elif o in ["-m", "--message"]:
        message=a
    elif o in ["-v", "--verbose"]:
        verbose=1
    #import pdb;pdb.set_trace()
    with Connection('amqp://'+username+':'+password+'@'+hostname+':'+port+'/') as conn:
        simple_queue = conn.SimpleQueue(queue)
        m_queue = '%s, sent at %s' % (message, datetime.datetime.today())
        simple_queue.put(m_queue)
        print('Sent: %s' % m_queue)
        simple_queue.close()

root@ctl3:~/bin# cat hello_con.py
#!/usr/bin/python
from kombu import Connection
import sys, os, getopt, re

def usage(message=None):
    print "Usage: %s [-h] [-u|--username <user_name>] [-w|--password <password>] [-h|--host <rabbitmq_hostname>] [-p|--port <service_port>] [-q|--queue <queue_name>]" % (sys.argv[0])
    print "\t-h|--help: show this message"
    print "\t-v|--verbose: include details in output"
    print "\t-u|--username: Username"
    print "\t-w|--password: User password"
    print "\t-h|--hostname: Rabbitmq hostname or IP"
    print "\t-p|--port: Rabbitmq service port"
    print "\t-q|--queue: Rabbitmq queue name"
    sys.exit(-1)

(opts, args) = getopt.getopt(sys.argv[1:], "?vu:wh:p:q:", ["Rabbitmq test", "help", "verbose", "username", "password", "host", "port", "queue", "message"])

for o, a in opts:
    if o in ["-?", "--help"]:
        usage()
    elif o in ["-u", "--username"]:
        username=a
    elif o in ["-w", "--password"]:
        password=a
    elif o in ["-h", "--host"]:
        host=a
    elif o in ["-p", "--port"]:
        port=a
    elif o in ["-q", "--queue"]:
        queue=a
    elif o in ["-v", "--verbose"]:
        verbose=1

    with Connection('amqp://'+username+':'+password+'@'+host+':'+port+'/') as conn:
        simple_queue = conn.SimpleQueue(queue)
        message = simple_queue.get(block=True, timeout=1)
        print("Received: %s" % message.payload)
        message.ack()
        simple_queue.close()

```

And to verify the rabbitmq deployment by publishing to and subscribing from the deployed rabbitmq.

```
# Publish 50 messages to the queue "test_queue" with 5 concurrently
#seq 50 | xargs -n 5 -P5 -I{} ./hello_pub.py -u guest -w guest -h 10.11.1.134 -p 5772 -q test_queue -m "from ctl3 number: {}"
```

```
#Subscribe the "test_queue" for 50 times with 5 concurrently
#seq 5 | xargs -n 5 -P5 -I{} ./hello_con.py -u guest -w guest -h 10.11.1.134 -p 5772 -q test_queue
```